

Locks vs Variables

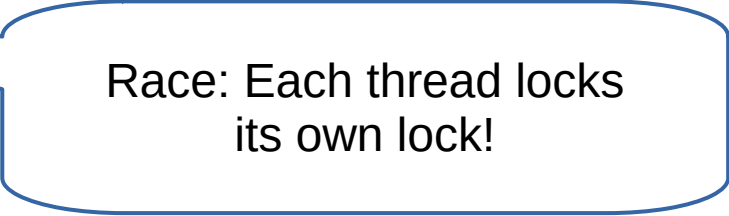
- “Which lock corresponds to each data item?”
- Multiple threads accessing some data item concurrently must have acquired the same lock
- Not automatic / not checked
- It is up to the developer to ensure this!

Pitfall: Automatic variables

- Variables in methods are created on every invocation
- The method might still access shared state:
 - Instance variables
 - Class (static) variables

Wrong

```
class SomeClass {  
    int s;  
    void doSomething() {  
        Lock l = new ReentrantLock();  
        l.lock();  
        s = s+1;  
        l.unlock();  
    }  
}
```



Race: Each thread locks
its own lock!

Solution

```
class SomeClass {  
    int s;  
    Lock l = new ReentrantLock();  
    void doSomething() {  
        l.lock();  
        s = s+1;  
        l.unlock();  
    }  
}
```

Solution: Same scope for
shared state and lock

Pitfall: Class/global variables

- Variables marked with “static” in Java are global and (probably) need concurrency control
 - Not if marked “final”
 - Not if the class is used by a single thread

Wrong

```
class SomeClass {  
    private static int s;  
    void doSomething() {  
        s = s+1;  
    }  
}
```



Race!

Still wrong

```
class SomeClass {  
    private Lock l = new ReentrantLock();  
    private static int s;  
    void doSomething() {  
        l.lock();  
        s = s+1;  
        l.unlock();  
    }  
}
```

There is one lock for each object, but s is shared!

Solution

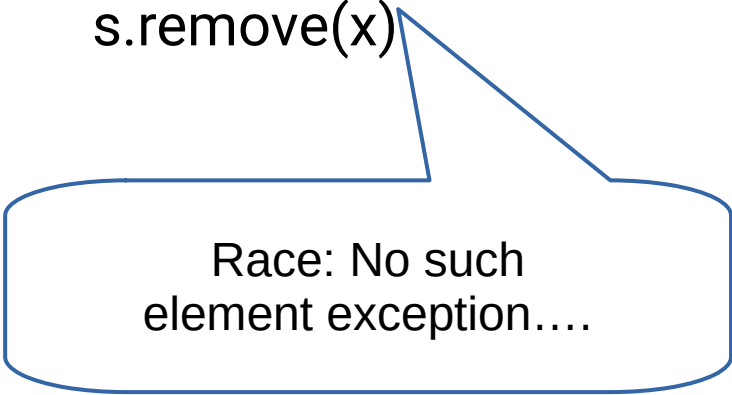
```
class SomeClass {  
    private static Lock l = new ReentrantLock();  
    private static int s;  
    void doSomething() {  
        l.lock();  
        s = s+1;  
        l.unlock();  
    }  
}
```


Pitfall: Encapsulated locks

- Keep variables and the corresponding lock encapsulated within the same object
- (The default using old-style “synchronized” in Java.)

Wrong

```
class SomeClass {  
    SomeState s;  
    void doSomething() {  
        if (s.contains(x))  
            s.remove(x)  
    }  
}
```



Race: No such
element exception....

```
class SomeState {  
    private Lock l;  
    boolean contains(...) {  
        l.lock(); ... l.unlock();  
    }  
    void remove(...) {  
        l.lock(); ... l.unlock();  
    }  
}
```

Solution

```
class SomeClass {  
    private Lock l;  
    SomeState s;  
    void doSomething() {  
        l.lock();  
        if (s.contains(x))  
            s.remove(x)  
        l.unlock();  
    }  
}
```

Now useless...

```
class SomeState {  
    private Lock l;  
    boolean contains(...) {  
        l.lock(); ... l.unlock();  
    }  
    void remove(...) {  
        l.lock(); ... l.unlock();  
    }  
}
```

Better solution

```
class SomeClass {  
    private Lock l;  
    SomeState s;  
    void doSomething() {  
        l.lock();  
        if (s.contains(x))  
            s.remove(x)  
        l.unlock();  
    }  
}
```

```
class SomeState {  
    private Lock l;  
    boolean contains(...) {  
        l.lock(); ... l.unlock();  
    }  
    void remove(...) {  
        l.lock(); ... l.unlock();  
    }  
}
```

Rely on locking by the callers. This is done by Java Collections (Lists, ...)


Pitfall: Shared vs thread-local state

- Program state often contains:
 - Local thread state in workers
 - Shared state, used by all threads
- Both are objects, with instance variables

Wrong

```
class Worker
    extends Thread {
    Lock l = new ...;
    SharedState s;
    void doSomething() {
        l.lock(); s.doit(); l.unlock();
    }
    public void run() {
        doSomething();
    }
}
```


```
class SharedState {
    public void doit() {
        ...
    }
}
```



Solution


```
class Worker
    extends Thread {
        SharedState s;
        void doit() {
            s.doit();
        }
        public void run() {
            doit();
        }
    }
```

```
class SharedState {
    Lock l = new ...;
    public void doit() {
        l.lock(); ... l.unlock();
    }
}
```



Quiz

- Two variables:
 - `int i=0, j=0;`
- Writer code:
 - `i=1; j=1;`
- Reader code:
 - `rj=j; ri=i; System.out.println(rj+", "+ri);`
- Possible results:
 - a) 0, 0 ✓
 - b) 1, 1 ✓
 - c) 0, 1 ✓
 - d) 1, 0 ✓




Why!?!?

} running
concurrently!

Memory order

- Steps to write a variable:
 1. Write to cache
 2. Flush cache to memory



(Image from <http://booksite.elsevier.com/9780123705914/?ISBN=9780123705914> . CC By-SA-3.0.)

Memory order

- Possible outcome with two variables:


1. Write i to cache

2. Write j to cache

3. Flush j from cache to memory

5. Flush i from cache to memory

4. Paradox observed if i,j read here!!



(Image from <http://booksite.elsevier.com/9780123705914/?ISBN=9780123705914> . CC By-SA-3.0.)


Pitfall: Read races

```
class X {  
    private Y y;  
    void changeY() {  
        l.lock();  
        tmp.i = 1;  
        y = tmp;  
        l.unlock();  
    }  
    int getY() {return y.i;}  
}
```

- Can we omit synchronization in getters?

Pitfall: Read races

```
class X {  
    private Y y;  
    void changeY() {  
        l.lock();  
        tmp.i = 1;  
  
        y = tmp;  
        l.unlock();  
    }  
    int getI() { return y.i;}  
}
```



- Can we omit synchronization in getters?
 - **NO!**
- Can read inconsistent Y fields!
- In this case:
 - reader might not see `y.i == 1!!!!`

Pitfall: Collections and getters

- Getter methods may return references to shared collections (or other mutable objects)
 - Iterators include references to the original object!

Wrong

```
class SomeClass {  
    private Lock l = new ReentrantLock();  
    private List l;  
    List getElements() {  
        try { l.lock();  
            return l;  
        } finally { l.unlock(); }  
    }  
}
```

```
SomeClass s = ...;  
List l = l.getElements();  
l.add(...);
```



Race!

Wrong

```
class SomeClass {  
    private Lock l = new ReentrantLock();  
    private List l;  
    Iterator getElements() {  
        try { l.lock();  
            return l.iterator();  
        } finally { l.unlock(); }  
    }  
}
```

```
SomeClass s = ...;  
    Iterator i = l.getElements();  
    while(i.hasNext())
```

...

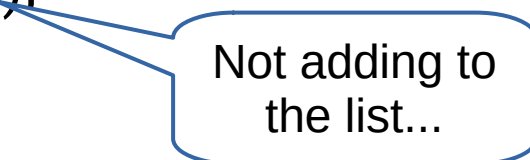


Race!

Still wrong

```
class SomeClass {  
    private Lock l = new ReentrantLock();  
    private List l;  
    List getElements() {  
        try { l.lock();  
            return l.clone();  
        } finally { l.unlock(); }  
    }  
}
```

```
SomeClass s = ...;  
List l = l.getElements();  
l.add(...);
```




Not adding to
the list...

...reconsider encapsulated lock!

Summary

- There is no simple rule to match locks with variables
- Some thinking needed... :-)



Scaling up

- Example:
 - In a distributed database table with millions of records
 - Executing “select sum(x) from ... where ...” queries
 - Updating records
- Do we use a single lock?
 - Cannot run more than one query at the same time
- Do we use a lock for each line?
 - Way too many individual locks!

Readers-Writers locks

- Strict mutual exclusion with locks is too conservative:
 - More than one reader would not be a problem
 - A writer must exclude all others (readers and writers)
- Different methods for readers and writers:

```
interface ReadWriteLock {  
    Lock readLock();  
    Lock writeLock();  
}
```
- More costly than a simple lock



Readers-Writers locks

- R/W locks also known as shared locks in database management systems:
 - Readers lock \Leftrightarrow Shared mode
 - Writers lock \Leftrightarrow eXclusive mode
- Behavior described by a compatibility matrix:

Mode	R/S	W/X
R/S	Yes	No
W/X	No	No

Readers-Writers example

```
int v;  
Lock l = new ReentrantLock();  
void doSomething() {  
    l.lock();  
    v++;  
    l.unlock();  
}  
int getV() {  
    try { l.lock();  
        return v;  
    } finally { l.unlock(); }  
}
```

```
int v;  
ReadWriteLock l = new ReentrantReadWriteLock();  
void doSomething() {  
    l.writeLock().lock();  
    v++;  
    l.writeLock().unlock();  
}  
int getV() {  
    try { l.readLock().lock();  
        return v;  
    } finally { l.readLock().unlock(); }  
}
```

... not worth it for such simple operations!

Revisiting collections with 2PL+RW lock

- ```
void shoot(String sn, String tn) {
 l.readLock().lock();
 Player s = players.get(sn);
 Player t = players.get(tn);
 Stream.of(sn,tn).sorted()
 .forEach(n→players.get(n).l.lock());
 l.readLock().unlock();
 t.life--;
 t.l.unlock();
 s.score++;
 s.l.unlock();
}
```

Allow multiple  
threads to acquire  
locks concurrently

Sorting is needed  
again

# Readers-Writers fairness

- Priority to readers
  - Allow more readers in, even if a writer is waiting
  - The writer may starve...
- Priority to writers
  - Do not allow more readers in, if a writer is waiting
  - Less concurrency among readers

# Lock managers

- Individual locks inefficient for huge collections of objects
  - A lock object uses memory even when not in use
- A lock manager provides locks on demand:

```
interface LockManager {
 void lock(Object name);
 void unlock(Object name);
}
```

lookup lock l for "name" in map  
if it doesn't exist:  
    create it and add to map  
l.lock()!

lookup lock for "name"  
l.unlock()  
if nobody else is using it:  
    remove it from map



# Lock managers

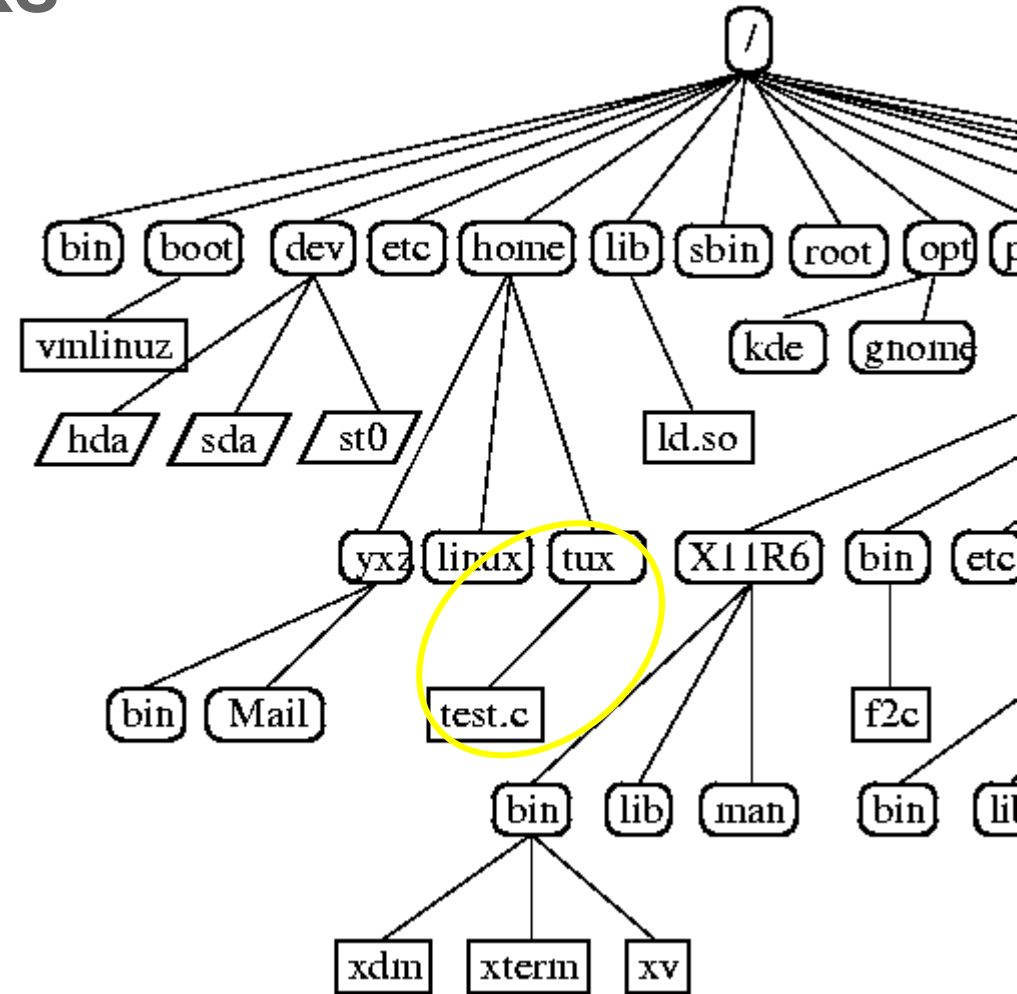
- Usually provides shared/exclusive semantics:

```
enum Mode { SHARED, EXCLUSIVE };
```

```
interface LockManager {
 void lock(Object name, Mode mode);
 void unlock(Object name);
}
```

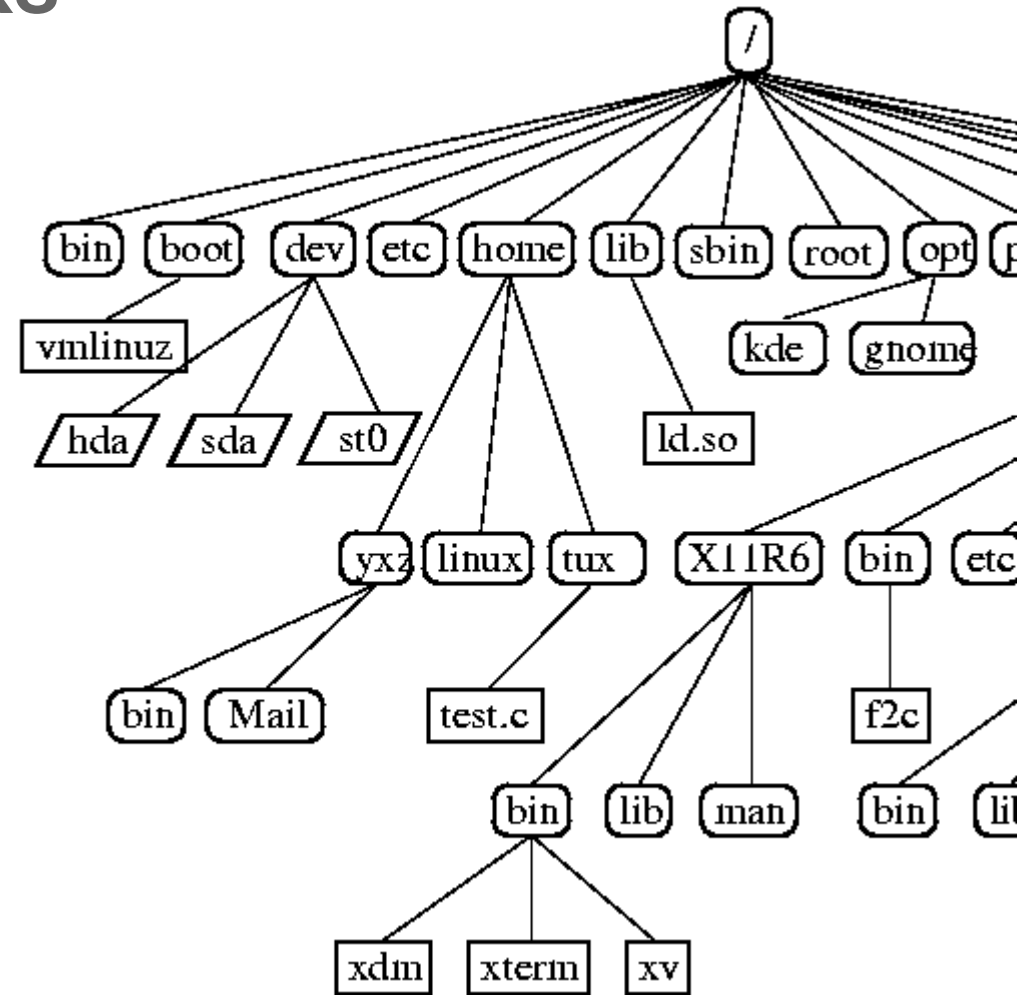
# Multiple granularity locks

- Motivation:
  - Locking `/home/tux/*`
  - Assume large number of files
- Inefficient even with a lock manager
- Idea: Take advantage of hierarchical namespace and lock folders



# Multiple granularity locks

- Protocol:
  - “intention” locks on containers
  - “actual” locks on the target
- Intention locks conflict with actual locks, not with other intention locks
- Combine with (S)hared and e(X)clusive semantics



# Compatibility matrix

- An MGL is defined by a compatibility matrix:

The diagram shows a 5x5 compatibility matrix. A blue bracket on the left, labeled 'acquired', spans the rows. A blue bracket on top, labeled 'requesting...', spans the columns. The matrix cells are colored: yellow for 'Yes' and red for 'No'.

| Mode | IS  | IX  | S   | X  |
|------|-----|-----|-----|----|
| IS   | Yes | Yes | Yes | No |
| IX   | Yes | Yes | No  | No |
| S    | Yes | No  | Yes | No |
| X    | No  | No  | No  | No |

# Multiple granularity locks

- Shared lock /home/tux
- Shared lock on /home/tux/test.c
- Exclusive lock on /boot
  - Shared lock on /boot/vmlinuz
    - IS on /boot/ conflicts with X
  - Exclusive lock on /home
    - X on /home conflicts with IS

