

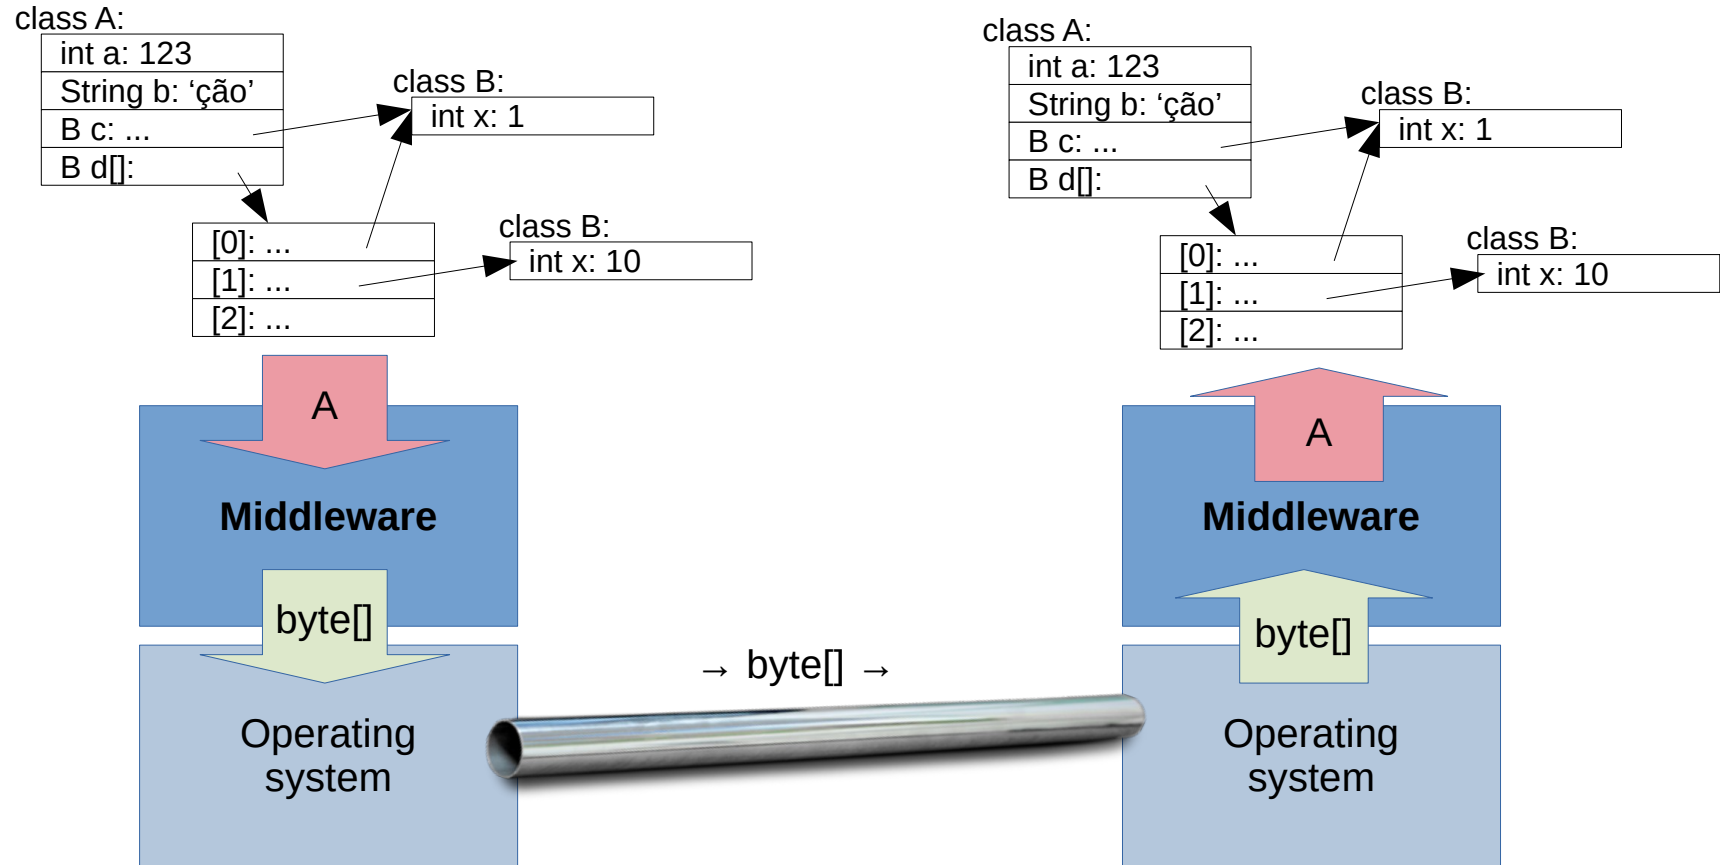
Sistemas Distribuídos

José Orlando Pereira

Departamento de Informática
Universidade do Minho



Serialization / Marshaling



Motivation

- Abstraction:
 - Messages as general purpose data structures
- Heterogeneity:
 - In space:
 - Different hardware
 - Different language / platform
 - Different middleware
 - In time:
 - Evolution of middleware and different versions co-existing in the same system

Roadmap

- Representation of basic data types
- Representation of composite data types
- Conversion code

Design issue: Text vs Binary


- Text formats:
 - Human readable and robust
 - Redundant and slower to parse
 - Examples: plain text, HTTP1.x, JSON
 - <https://json.org/example.html>
- Binary formats:
 - Compact and efficient
 - Opaque (harder to debug) and brittle
 - Examples: Java Data*Stream streams

Binary formats: Representation

- Endianness
 - An integer: 3735928559 / 0xDEADBEEF
 - Big endian bytes: { 0xDE, 0xAD, 0xBE, 0xEF }
 - Little endian bytes: { 0xEF, 0xBE, 0xAD, 0xDE }
- Character encoding
 - A string: “ção”
 - UTF8: { 0xC3, 0xA7, 0xC3, 0xA3, 0x6F }
 - Latin1: { 0xE3, 0xE7, 0x6F }

Binary formats: Alignement and padding

- Memory is addressed at byte offsets
- But accessed as multi-byte words
- Unaligned accesses are:
 - Slower; or
 - not allowed



- Binary formats may use padding for alignement

Binary formats: Example

- Example in Java:
 - java.io.DataOutput/DataInput

```
os.writeInt(123);  
os.writeUTF("çãõ");  
...
```
- Uses a common representation
 - Big endian
 - Modified UTF8 strings
 - IEEE standard floating point

Design issue: Implicit vs Explicit

- Explicit formats describe their own content (types and/or item names):

`<file>`

`<format>mp3</format>`

`<tags><tag>jazz</tag><tag>modern</tag></tags>`

`<size>5443236</size>`

`</file>`

- Implicit formats depend on custom code to read them
`mp3\0\0x2jazz\0modern\0\0x31\0x24...`

Design issue: Single or Multiple

- Agree on a common representation (aka “network byte order”):
 - Convert when sending
 - Convert when receiving
 - Even if sender and receiver are identical!
- Use sender representation:
 - Send with a tag
 - Depending on tag, convert when receiving

Composite types

- Non-contiguous in memory
 - Lists, trees, ...
- Contain additional information, not needed or meaningless over the network
 - Pointers, locks, ...

Composite types

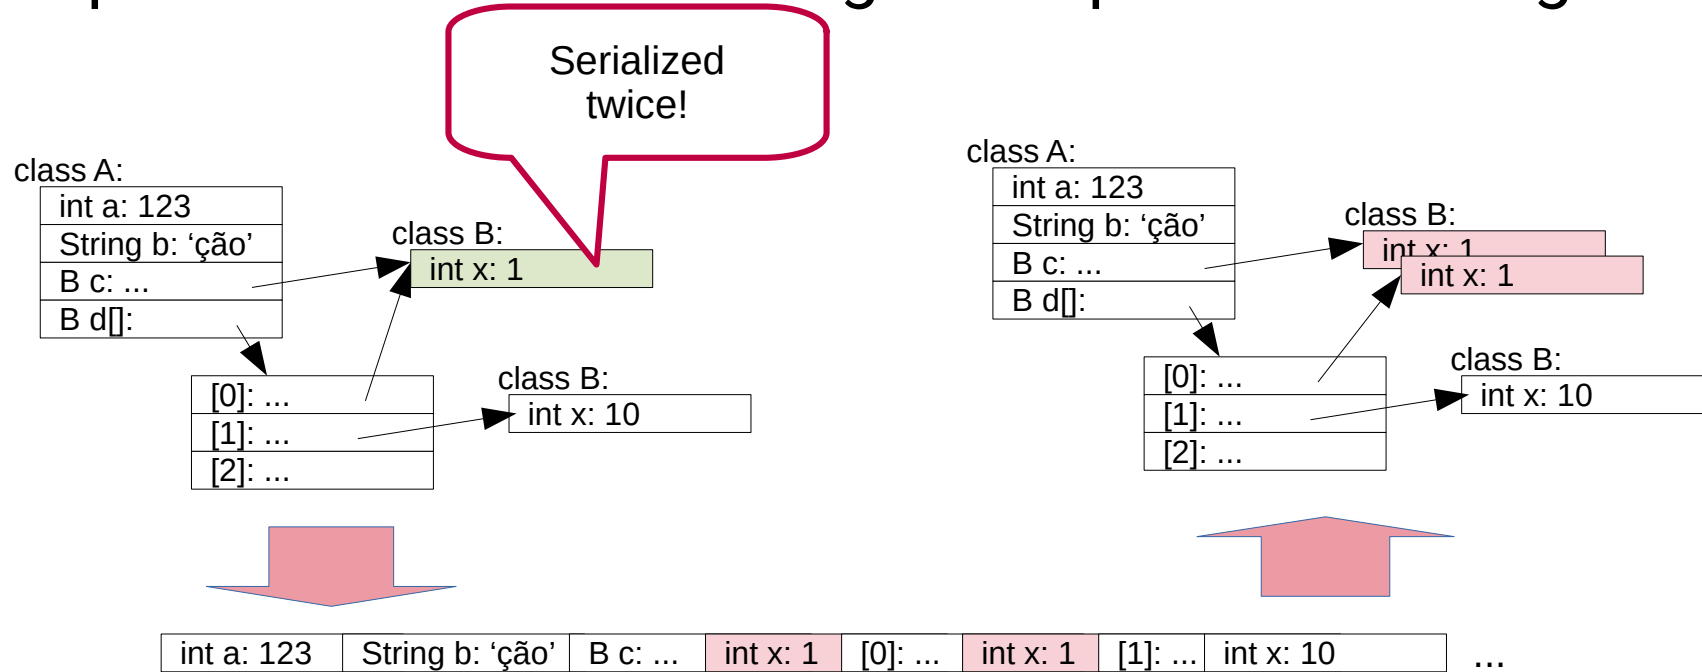
- Records
 - Enumerate each of the components
 - Include optional padding
- Objects (with subclassing) and unions:
 - Prefix content with a tag that identifies the actual option used
 - Use the tag to determine what to deserialize
- Optional items (e.g., nullable fields)
 - Prefix with a boolean indicating if present

Composite types: Collections

- Arrays, lists, sets, and maps
- First option:
 - Prefix with the number of components, then each of the components
 - Common in binary representations
 - Better if the size can be determined easily
- Second option:
 - Each of the components, then a special terminator value
 - Common in text-based representations
 - Better if the data structure can grow dynamically

Composite types: Graphs

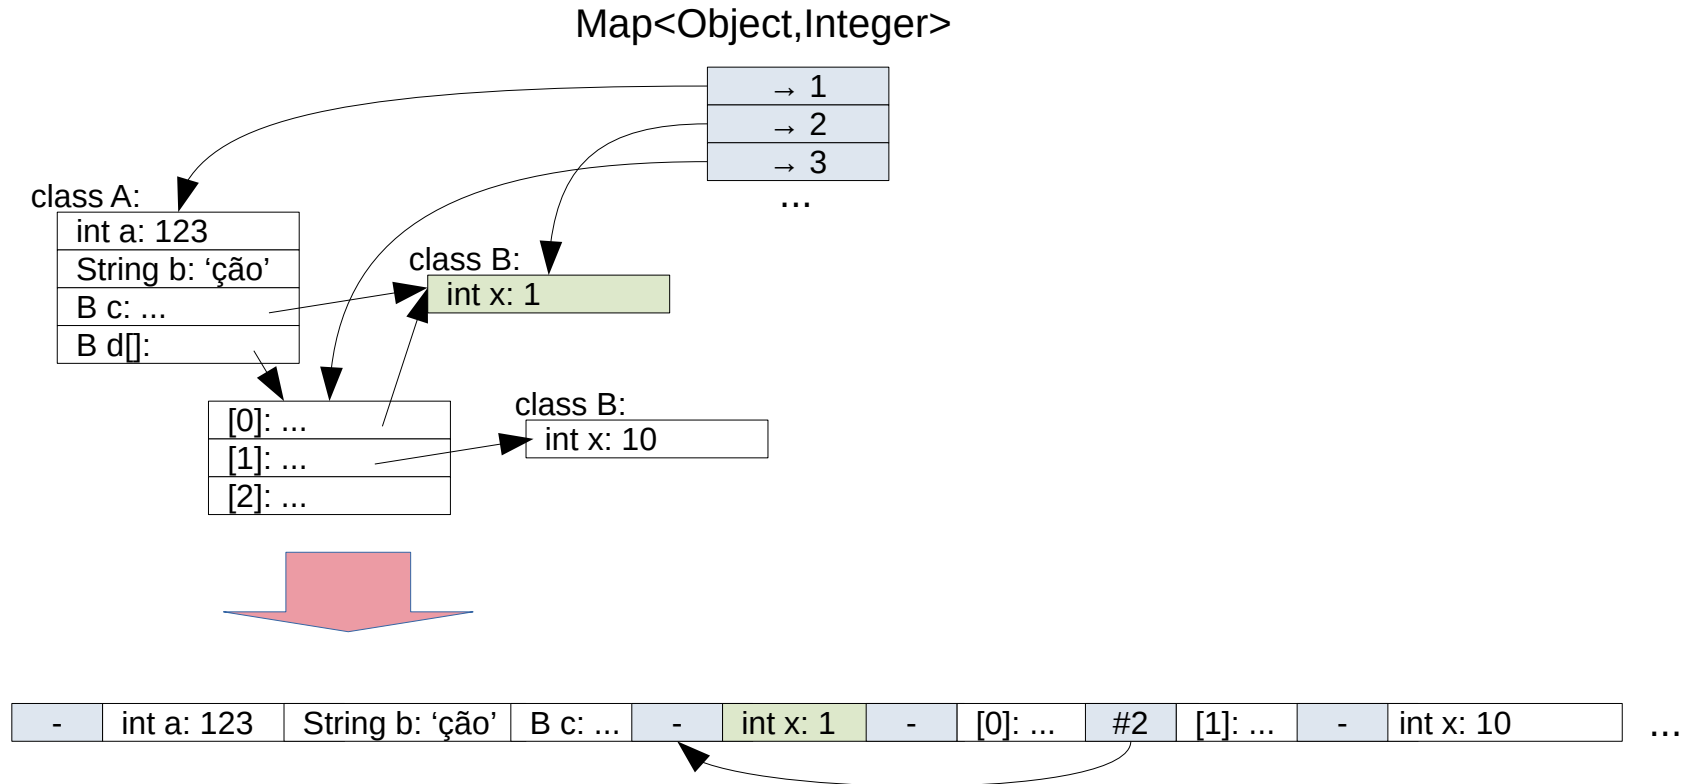
- Simple traversal is not enough with pointer aliasing:



- In fact, with cycles, simple traversal does not terminate (and generates unbounded data...)

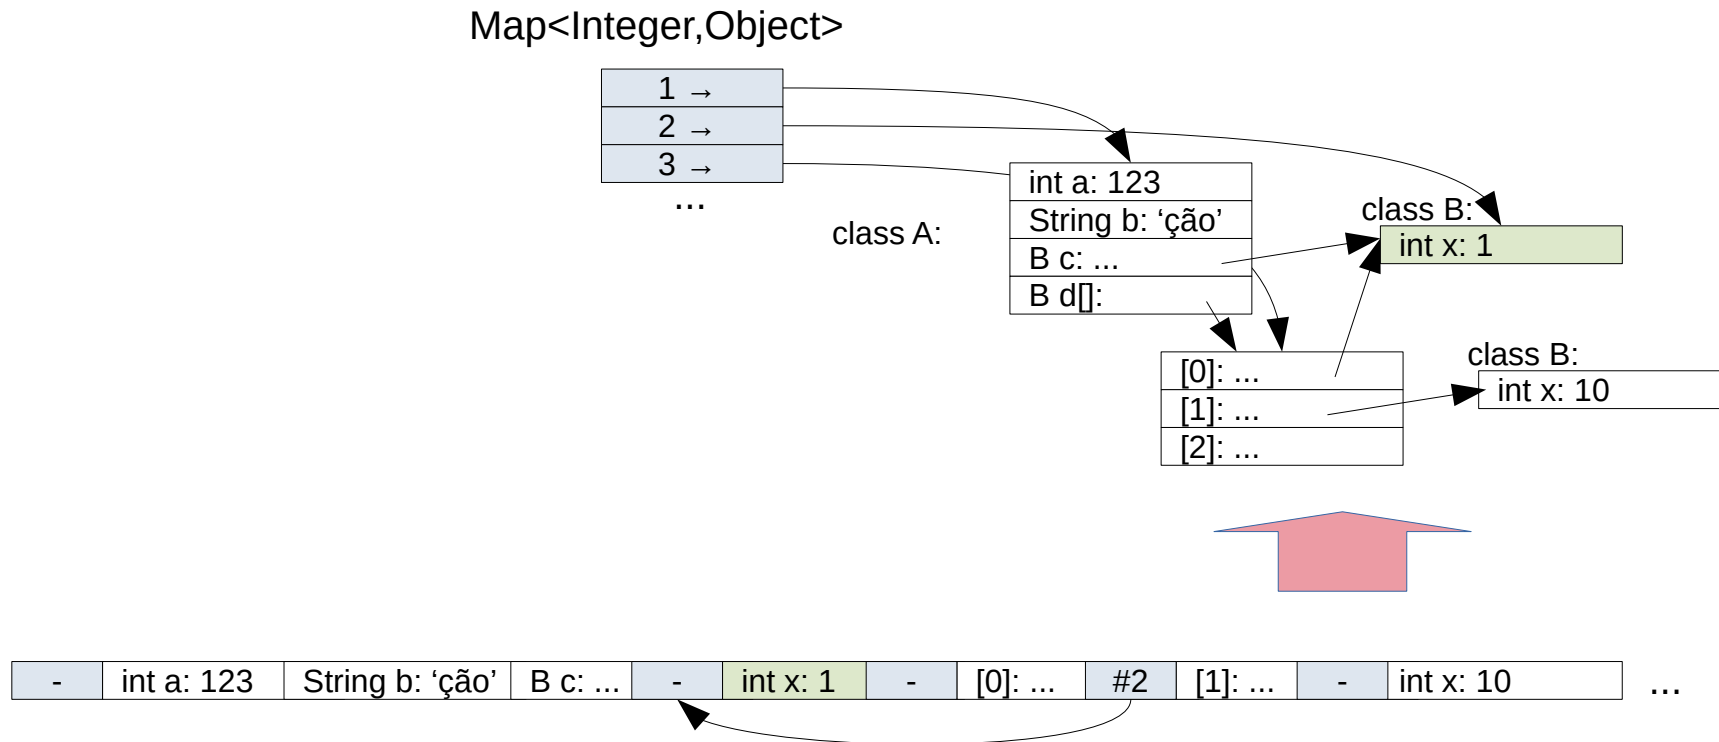
Composite types: Graphs

- Use tags and an auxiliary map while serializing:



Composite types: Graphs

- When deserializing, keep track of objects to restore pointers:



Composite types and graphs

- Example in Java:
 - `java.io.ObjectOutput/ObjectInput`
 - Uses `DataOutput/DataInput` for basic types
 - Recreates object graphs
- Very inefficient...

Design issue: Program vs Data first

- Program first: Data format inferred from an existing program
 - Convenient for development
 - Tied to a single language
 - Example: Java Object*Stream
- Data first: Program generated from an abstract description of data
 - Language and middleware independence
 - Forces developer to use new tools
 - Example: Protobuf

Conversion code

- Enumerate components to be written / read by recursively traversing data structures
- Filter methods
 - Tedious and error-prone
- Reflection
 - Allow overriding for transient data (Locks, caches, ...)
- Generated code
 - Interface language and compiler

Design issue: Versioning

- If the data structure changes, the receiver will break / decode corrupt data
- Allow a data structure to be modified with new versions of the program
- Make individual items optional and/or provide defaults
 - Example: Protobuf
- Allow versioning of the data structure as a whole
 - Example: Java Object*Stream

Design issue: Streaming vs Object model

- Streaming:
 - Data directly copied from / to external representation
 - Internal and external layouts exactly the same
 - Better efficiency: data copied only once
- Object model:
 - An intermediate model is built in memory
 - The model can be queried and traversed
 - Some applications might actually use the intermediate model directly

Common alternatives

- Simple text lines
- **Java Data*Stream**
- Java Serialization (Object*Stream)
- Kryo
- Protobuf
- JSON
- XML

Summary

- Heterogeneity as a key property of distributed systems
- Standard data representation and serialization/marshaling are needed to address it
- Tedious and error prone programming task that should be automated
- Key issue in access transparency