

Sistemas Distribuídos


José Orlando Pereira

Departamento de Informática
Universidade do Minho

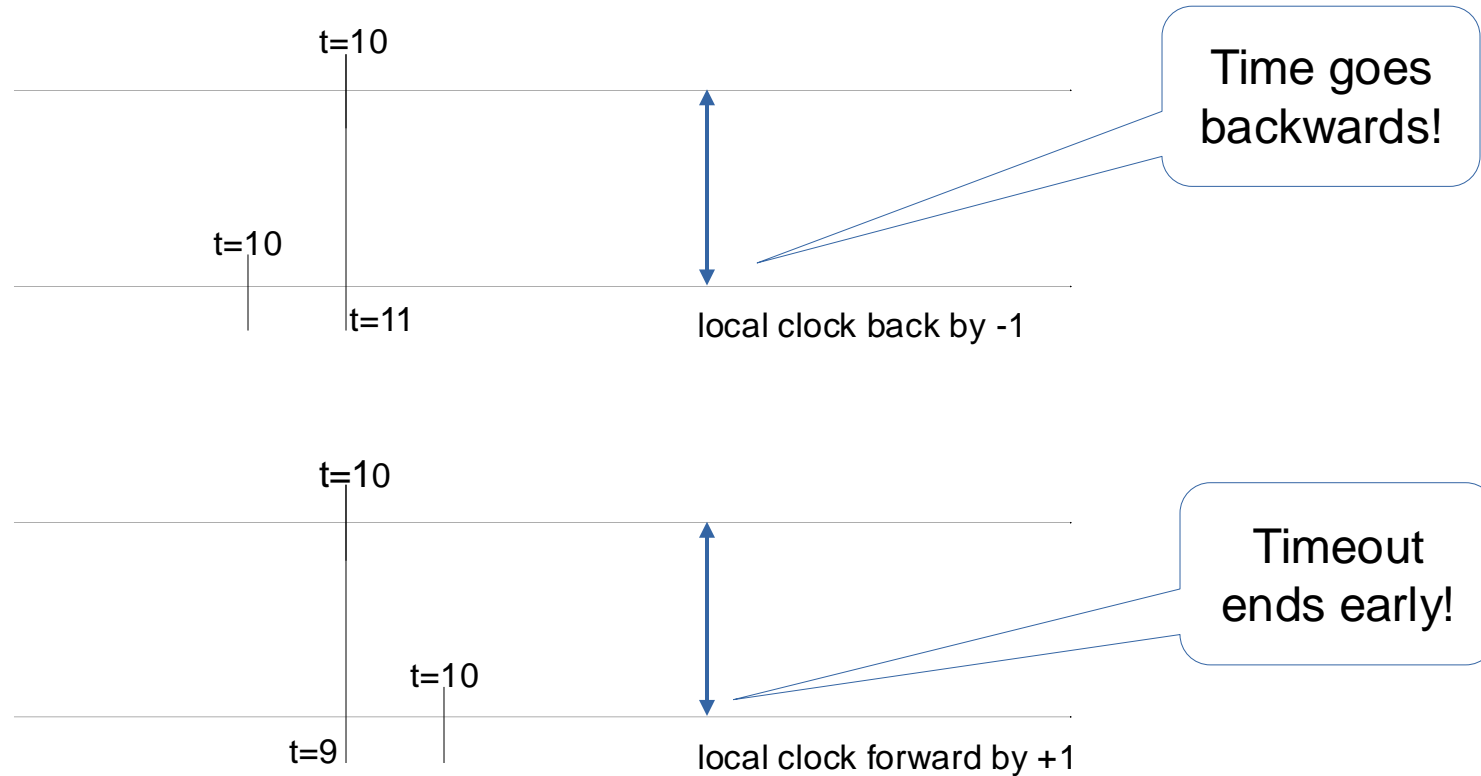


Clock synchronization

- Hardware clocks are not perfect and drift over time
- Clock skew can be a problem with:
 - shared files
 - certain algorithms...
- Ideally:




Clock synchronization



- The clock must be adjusted in small increments, over a longer period of time by making it faster or slower

Clock synchronization


- In practice, there are unpredictable
 - transmission delays
 - processing delays



- The best we can do is adjust by $(T2-T1)$ – (estimated message delay)

Clock synchronization


- What is the message delay?



- Network Time Protocol (NTP):
 - assume delays are the same = $((T2-T1)+(T4-T3))/2$
 - repeat several times and pick the smallest delay

Clock synchronization

- Reference Base Synchronization (RBS):
 - assume true broadcast medium (aprox. zero skew)




Mutual exclusion

- Solution in a distributed system?
- Recall the definition of mutual exclusion:
 - No two threads in the critical section
 - No deadlock / no starvation
- Comparison criteria:
 - Number of message hops to entering critical section
 - Load balancing

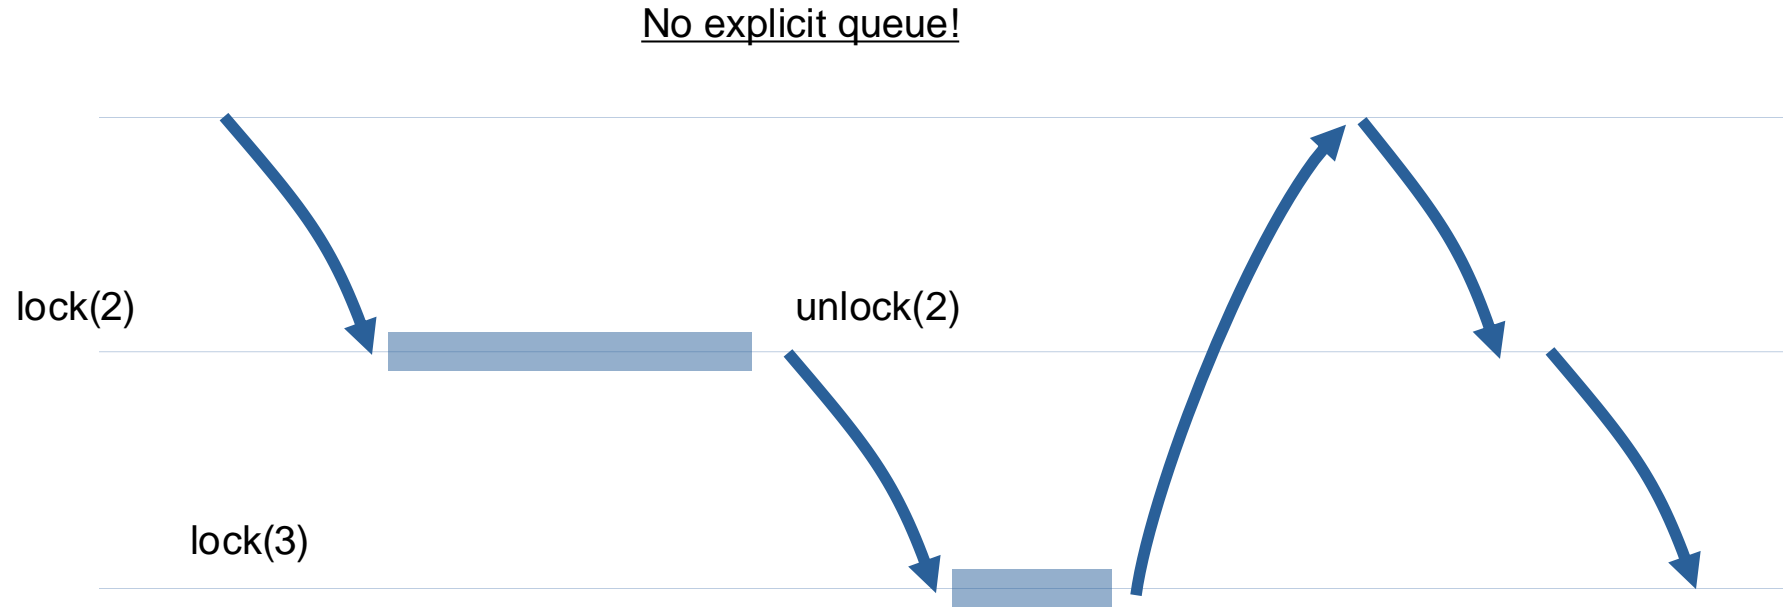
Mutual exclusion

- Centralized queue kept by a coordinator:
 - 1 round-trip to enter / asymmetric load




Mutual exclusion

- Exchange a token in a ring:
 - $n/2$ hops to entering / symmetric load




Mutual exclusion

- A distributed algorithm is hard to achieve:
 - As concurrent lock requests are received by different destinations in different orders, safety is not ensured




Mutual exclusion

- Taking advantage of synchronized clocks:
 - Assume $\delta > (\text{transmission delay} + \text{skew})$
 - Consider only messages up to $t - \delta$, order by timestamp




Timeouts

- Used to assess status of remote processes
- Tight timeouts are dangerous:
 - E.g., proportional to mean delay
 - Means low coverage
- Large timeouts are not useful:
 - E.g., proportional to high percentile
 - Taking advantage of time causes a very large performance penalty



Using real time

- Solutions that do not use time are more robust:
 - In wide area networks
 - With performance perturbations
- Solutions that do not use time might have better performance:
 - Run time proportional to mean delay
 - Even if more message exchanges are necessary




Asynchronous system model

- Assume no global time reference
- Assume no bounds on:
 - clock drift
 - processing time
 - message passing time
- Can we still solve important problems?

Time and causality

- What is special about time that makes it useful for distributed algorithms?




Time and causality

- $Clock(i)$ the time at which i happened
- If i precedes j then $Clock(i) < Clock(j)$
- For some event j :
 - When we are sure that there is no unknown i such that $Clock(i) < Clock(j)$
 - Then there is no i such that i precedes j
- **Can we build a logical clock with the same property?**

Lamport's logical clocks

- Local events: increment counter
- Send events: increment and then tag with counter
- Receive events: update local counter to maximum and then increment




Mutual exclusion

- Algorithm sketch:
 - Start by assuming that processes are continually exchanging messages over FIFO channels...
 - $ri[j]$ latest timestamp from j at i
 - Consider requests with $t \leq \min(ri[j], \text{for all } j)$
 - (akin to $t - \delta$!)
 - Order by timestamp, break ties by process id
- (The complete version is the Ricart-Agrawala distributed mutex algorithm)

Mutual exclusion

- 1 hop to enter / symmetric load



Conclusion

- The same approach used for the waiting queue in the mutex can be used for other deterministic applications
 - Replicated State Machine (RSM)
- Logical time is widely applicable in distributed systems to solve many problems